


# DSG-SOA-M 2024: - Linux Containers -

Dr. Andreas Schönberger

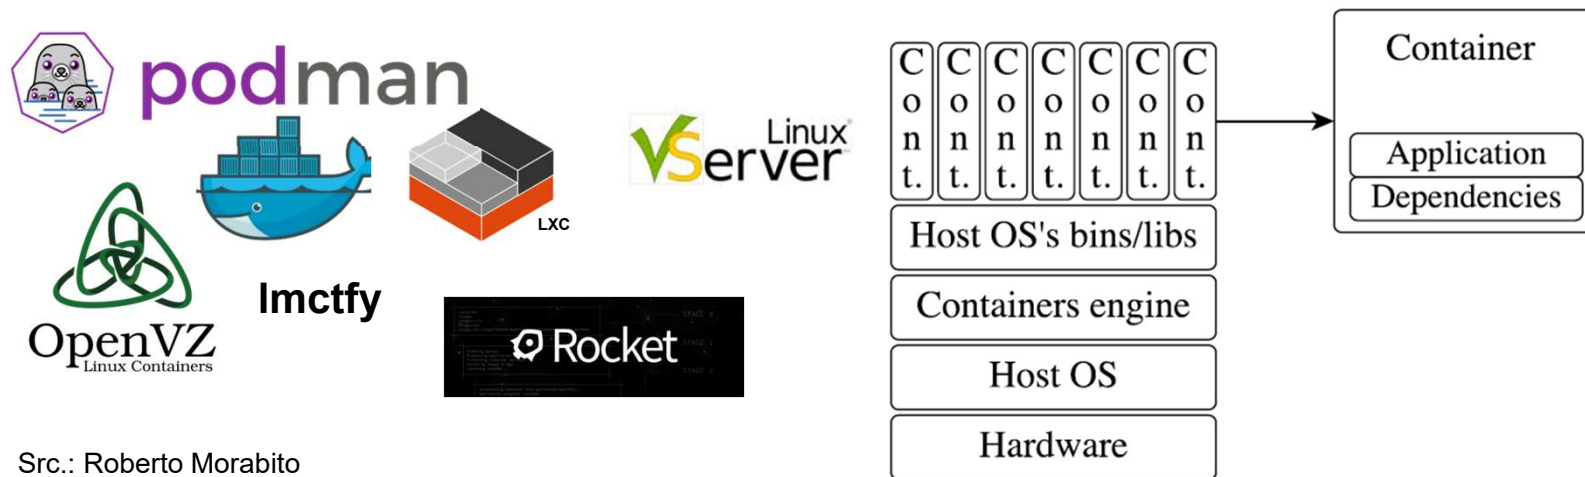
Distributed Systems Group  
Faculty Information Systems and Applied Computer Science  
University of Bamberg





# What are Linux Containers?


- ❑ Lightweight alternative to hypervisor-based virtualization
  - Higher density of virtualized instances
  - Efficient use of scarce embedded resources
- ❑ Containers implement isolation of processes at the OS level
- ❑ Run on top of the same shared OS kernel of the underlying host machine



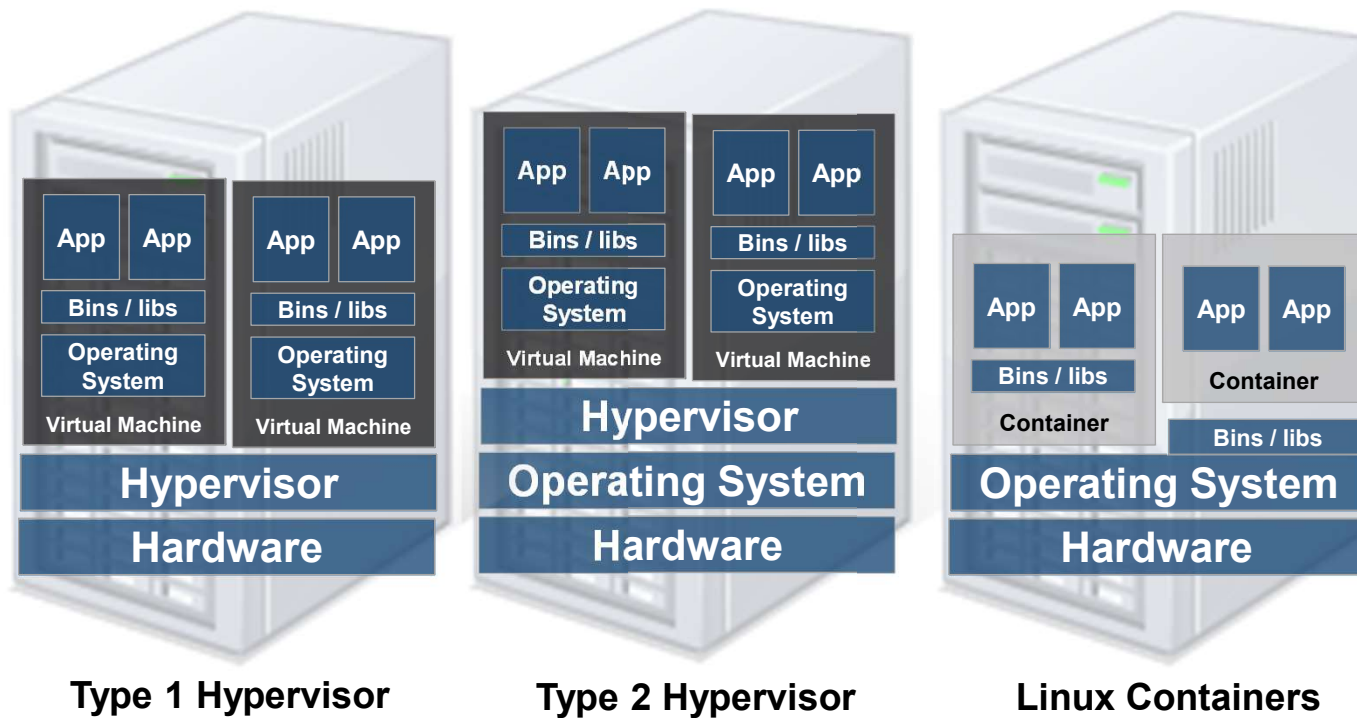
Src.: Roberto Morabito  
Ericsson Research



# Hypervisors vs. Linux Containers




Containers share the OS kernel of the host and thus are lightweight. However, each container must have the same OS kernel.



Src.:Boden Russell,  
IBM





# Where Do We Come from?


- ❑ Unix chroot (1979)
- ❑ FreeBSD Jails (2000)
- ❑ Linux VServer (2001)
- ❑ Solaris zones (2004)
- ❑ OpenVZ (2005)
- ❑ Linux Containers – LXC (2009)
- ❑ Docker (2013)
- ❑ Podman (2018)

Note: LXC as a technology != LXC tools




# Why Linux Containers?

---




Run many applications in a fast, portable and isolated way in many different environments!

Containers are like light-weight VMs.  
Small footprint, fast boot, shared kernel.




# Why Linux Containers – in more detail

- ❑ Provision in seconds / milliseconds
- ❑ Near bare metal runtime performance
- ❑ VM-like agility – it's still “virtualization”
- ❑ Flexibility
  - Containerize a “system”
  - Containerize “application(s)”
- ❑ Lightweight
  - Just enough Operating System
  - Minimal per container penalty
- ❑ Open source – free – lower TCO
- ❑ Supported with OOTB modern Linux kernel





Src.:Boden Russell,  
IBM

**Google trends - */xc***



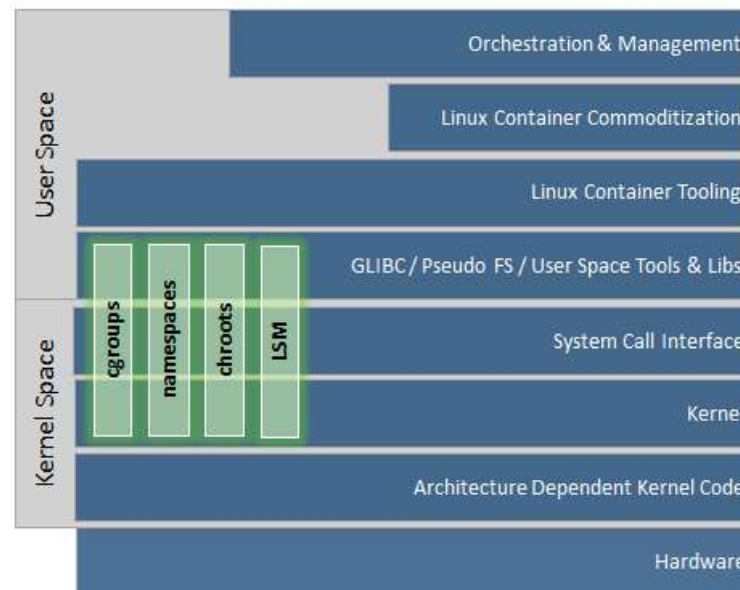
**Google trends - *docker***





# Ok, What is the Stack Beneath?


- ❑ LXC's are built on modern kernel features
  - cgroups; limits, prioritization, accounting & control
  - namespaces; process based resource isolation
  - chroot; apparent root FS directory
  - Linux Security Modules (LSM); Mandatory Access Control (MAC)
- ❑ User space interfaces for kernel functions
- ❑ LXC tools
  - Tools to isolate process(es) virtualizing kernel resources
- ❑ LXC commoditization
  - Dead easy LXC
  - LXC virtualization
- ❑ Orchestration & management
  - Scheduling across multiple hosts
  - Monitoring
  - Uptime



Src.:Boden Russell,  
IBM



# Linux cgroups – Core Tool for Resource Isolation



## □ Functionality

- Access; which devices can be used per cgroup
- Resource limiting; memory, CPU, device accessibility, block I/O, etc.
- Prioritization; who gets more of the CPU, memory, etc.
- Accounting; resource usage per cgroup
- Control; freezing & check pointing
- Injection; packet tagging


## □ Usage

- cgroup functionality exposed as “resource controllers” (aka “subsystems”)
- Subsystems mounted on FS
- Top-level subsystem mount is the root cgroup; all procs on host
- Directories under top-level mounts created per cgroup
- Procs put in tasks file for group assignment
- Interface via read / write pseudo files in group

Src.:Boden Russell,  
IBM








# Linux cgroups Subsystems

- ❑ cgroups provided via kernel modules (should be 3.10 or higher)
  - Not always loaded / provided by default
  - Locate and load with modprobe
- ❑ See: <https://www.kernel.org/doc/Documentation/cgroup-v1/>

Subsystem	Tunable Parameters
<b>blkio</b>	<ul style="list-style-type: none"><li>- Weighted proportional block I/O access. Group wide or per device.</li><li>- Per device hard limits on block I/O read/write specified as bytes per second or IOPS per second.</li></ul>
<b>cpu</b>	<ul style="list-style-type: none"><li>- Time period (microseconds per second) a group should have CPU access.</li><li>- Group wide upper limit on CPU time per second.</li><li>- Weighted proportional value of relative CPU time for a group.</li></ul>
<b>cpuset</b>	<ul style="list-style-type: none"><li>- CPUs (cores) the group can access.</li><li>- Memory nodes the group can access and migrate ability.</li><li>- Memory hardwall, pressure, spread, etc.</li></ul>
<b>devices</b>	<ul style="list-style-type: none"><li>- Define which devices and access type a group can use.</li></ul>
<b>freezer</b>	<ul style="list-style-type: none"><li>- Suspend/resume group tasks.</li></ul>
<b>memory</b>	<ul style="list-style-type: none"><li>- Max memory limits for the group (in bytes).</li><li>- Memory swappiness, OOM control, hierarchy, etc..</li></ul>
<b>hugetlb</b>	<ul style="list-style-type: none"><li>- Limit HugeTLB size usage.</li><li>- Per cgroup HugeTLB metrics.</li></ul>
<b>net_cls</b>	<ul style="list-style-type: none"><li>- Tag network packets with a class ID.</li><li>- Use tc to prioritize tagged packets.</li></ul>
<b>net_prio</b>	<ul style="list-style-type: none"><li>- Weighted proportional priority on egress traffic (per interface).</li></ul>



# How to Use cgroups - Linux cgroups Pseudo FS Interface



- ❑ Linux pseudo FS is the interface to cgroups
  - Read / write to pseudo file(s) in your cgroup directory
- ❑ Some libs exist to interface with pseudo FS programmatically


Find details  
on [kernel.org](https://kernel.org)

/sys/fs/cgroup/my-lxc


```
-- blkio
|-- blkio.io_merged
|-- blkio.io_queued
|-- blkio.io_service_bytes
|-- blkio.io_serviced
|-- blkio.io_service_time
|-- blkio.io_wait_time
|-- blkio.reset_stats
|-- blkio.sectors
|-- blkio.throttle.io_service_bytes
|-- blkio.throttle.io_serviced
|-- blkio.throttle.read_bps_device
|-- blkio.throttle.read_iops_device
|-- blkio.throttle.write_bps_device
|-- blkio.throttle.write_iops_device
|-- blkio.time
|-- blkio.weight
|-- blkio.weight_device
|-- cgroup.clone_children
|-- cgroup.event_control
|-- cgroup.procs
|-- notify_on_release
|-- release_agent
|-- tasks
-- cpu
-- ...
-- ...
-- perf_event
```

echo "8:16 1048576" >  
blkio.throttle.read\_bps\_de  
vice

cat blkio.weight\_device  
dev weight  
8:1 200  
8:16 500




Src.:Boden Russell,  
IBM




# Linux cgroups: CPU Usage

- ❑ Use CPU shares (and other controls) to prioritize jobs / containers
- ❑ Carry out complex scheduling schemes
- ❑ Segment host resources
- ❑ Adhere to SLAs




Src.:Boden Russell,  
IBM



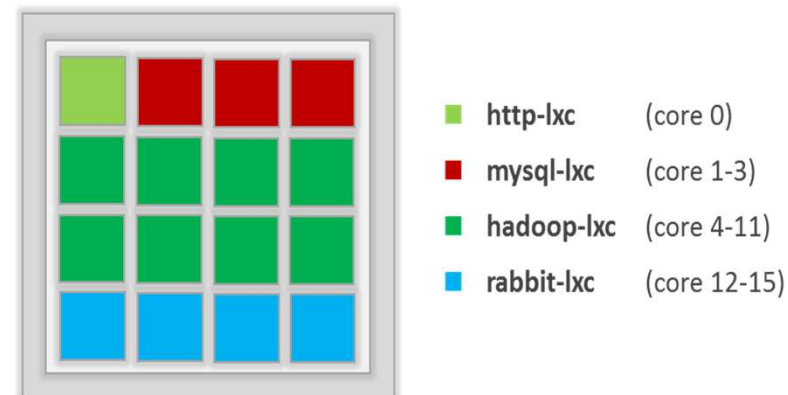
# Linux cgroups: CPU Pinning

- ❑ Pin containers / jobs to CPU cores
- ❑ Carry out complex scheduling schemes
- ❑ Reduce core switching costs
- ❑ Adhere to SLAs


Without CPU Core Pinning



With CPU Core Pinning

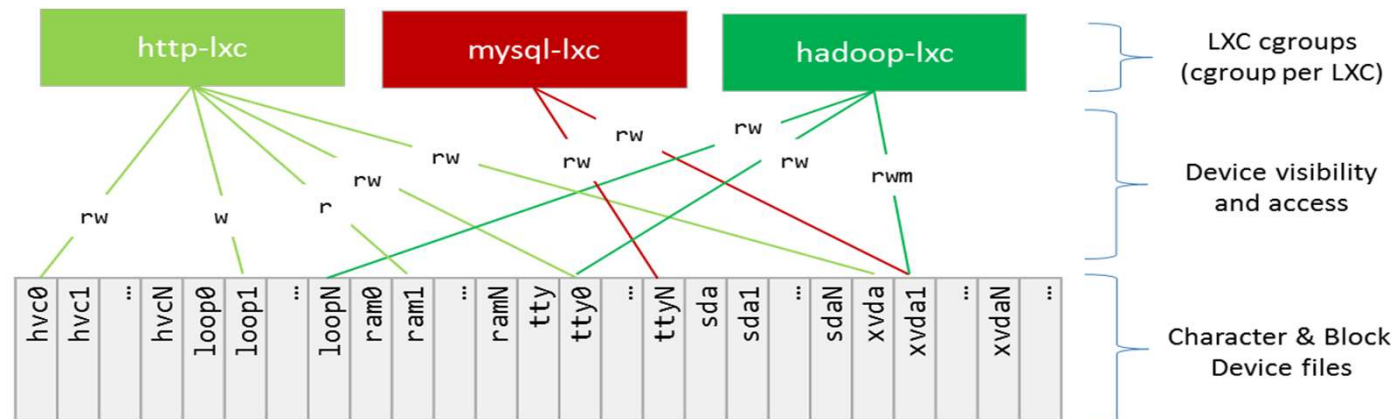


Src.:Boden Russell,  
IBM




# Linux cgroups: Device Access

- ❑ Limit device visibility; isolation
- ❑ Implement device access controls
  - Secure sharing
- ❑ Segment device access
- ❑ Device whitelist / blacklist



Src.:Boden Russell,  
IBM



# Linux Namespaces

## ❑ Why namespaces?

Namespaces provide isolation of Linux resources for a group of processes.

Src.:Boden Russell,  
IBM

## ❑ Functionality

### ▪ Provide process level isolation of global resources

- **MNT** (mount points, file systems, etc.)
- **PID** (process)
- **NET** (NICs, routing, etc.)
- **IPC** (System V IPC resources)
- **UTS** (host & domain name)
- **USER** (UID + GID)

### ▪ Process(es) in namespace have illusion they are the only processes on the system

### ▪ Generally constructs exist to permit “connectivity” with parent namespace

## ❑ Usage

### ▪ Construct namespace(s) of desired type


### ▪ Create process(es) in namespace (typically done when creating namespace)

### ▪ If necessary, initialize “connectivity” to parent namespace

### ▪ Process(es) in name space internally function as if they are only proc(s) on system




# Typical Container Engine / Tooling Functionality



- ❑ Standard container operations to realize / manage LXC's
  - Run, start, stop, delete, attach, snapshot, etc.
- ❑ APIs
  - CLI, Web API (RESTful or other), Automation files (e.g. dockerfile)
- ❑ Images
  - Format (typically RAW), required dev files, etc..
  - Image repository or catalog
  - Container FS layers (union FS style)
  - Import / export images and containers
- ❑ Metrics
  - CPU, memory, block IO, etc..
- ❑ Eventing
  - Push events on container changes
- ❑ Logs
  - Inspect / manage container logs
- ❑ Checkpointing
  - Freeze a container's state

Src.:Boden Russell,  
IBM





# What are / were the Liabilities?

- ❑ Live migration still an excitement
- ❑ Full orchestration across resources  
(compute / storage / networking)
- ❑ Fears of security
- ❑ Not a well known technology... but maturing very fast
- ❑ Integration with existing virtualization and Cloud tooling
- ❑ Not much / any industry standards
- ❑ Missing skillset
- ❑ Slower upstream support due to kernel dev process

Src.:Boden Russell,  
IBM

