

DSG-SOA-M 2024: - Docker -

Dr. Andreas Schönberger,
Stefan Kolb, Marcel Großmann,
Johannes Manner

Lehrstuhl für Praktische Informatik
Fakultät WIAI
Otto-Friedrich-Universität Bamberg



Docker is...

- ❑ name of the open source project Docker
- ❑ name of the company Docker, Inc., core sponsor of the Docker project
- ❑ founded in 2013 and has attracted significant attention since then




Figure from 03/2018

<https://blog.docker.com/2018/03/5-years-later-docker-journey/>

- ❑ a gartner report states: „By 2023, more than 70% of global organizations will be running more than two containerized applications in production, up from less than 20% in 2019.”

<https://www.itopstimes.com/contain/enterprise-container-strategy-its-time-to-jump-on-board/>



□ More statistics . . .






Figure from 07/2020

<https://www.docker.com/blog/docker-index-dramatic-growth-in-docker-usage-affirms-the-continued-rising-power-of-developers/>

- 242B Pulls in July 2020 compared to 37B in March 2018 (654% growth and still ongoing)



What Percentage of Apps are Containerized Today?



<https://www.stackrox.com/post/2020/03/6-container-adoption-trends-of-2020/>




What Percentage of Containerized Apps run in Production? (2020)



<https://www.stackrox.com/post/2020/03/6-container-adoption-trends-of-2020/>




What Percentage of Containerized Apps run in Production? (2021)




What percentage of your containerized apps are running in production?

Spring 2019
Winter 2020
Fall 2020



<https://www.stackrox.com/kubernetes-adoption-security-and-market-share-for-containers/>




The Docker Hype (is it real?) – some remarks

- ❑ Larger companies are the Early Adopters
(datadog 2016)
- ❑ 2/3 of companies that try Docker adopt It
(datadog 2016)
- ❑ There are 460,000 Dockerized applications, a 3100% growth over two years. Over 4 billion containers have been pulled so far.
(Docker's CEO Ben Golub at DockerCon 2016)
- ❑ Real Docker adoption Is Up 40% in One Year
(datadog 2017)
- ❑ Larger companies are still leading adoption
(datadog 2017)
- ❑ Docker now runs on 15% of the hosts we (datadog) monitor
(datadog 2017)
- ❑ “Docker reaching 54 percent adoption among larger companies”
(RightScale 2018 State of the Cloud Report)
- ❑ Docker now runs on more than 20% of the hosts we (datadog) monitor
(datadog 2018)

2020: Ok, let's stop asking the “Is it real?” question





Why all the Buzz?


- ❑ Docker uses “old” concepts, as used in, e.g., BSD Jails, Solaris Zones
- ❑ but implements them with modern technologies in smart ways (?)

“Smart ways” of Docker

- **Make container technology very accessible**
- **Foster community for collaboration** (Open Source, Meetups, Conventions, ...)
- **Offer huge ecosystem**

⇒ **Disruptive technology (?)**






What does Docker do?

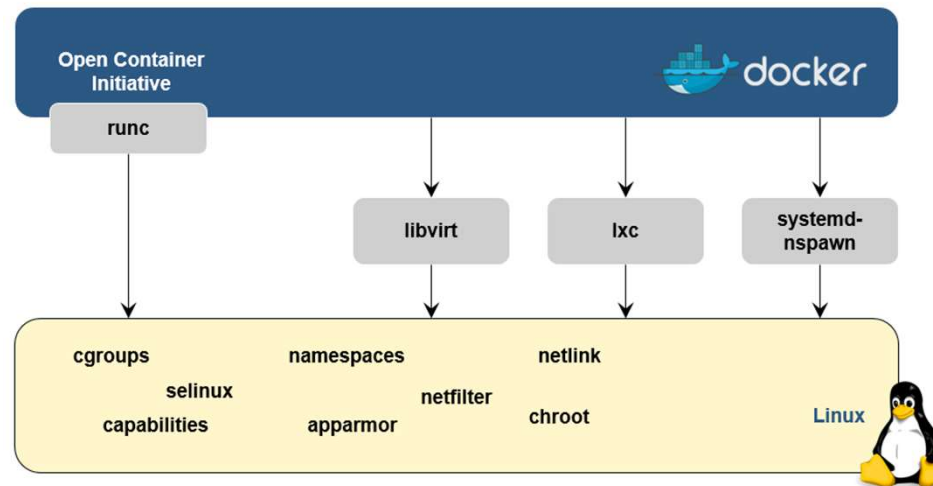
- ❑ “Docker allows you to package an application with all of its dependencies into a **standardized** unit for software development.”
- ❑ “Docker containers wrap up a piece of software in a complete filesystem that **contains everything** it needs to run: code, runtime, system tools, system libraries – anything you can install on a server. This guarantees that it will **always run the same**, almost regardless of the environment it is running in.”


⇒ **Build, ship and run everywhere!**



The Underlying Technologies

- Docker makes Linux containers usable in a simple fashion (so Docker builds upon Linux containers, some windows integration is offered)
- As a recap: containers share the host's kernel! (cgroups, namespaces)
- “The **Open Container Initiative** is an open governance structure for the express purpose of creating open industry standards around container formats and runtimes.” (<https://opencontainers.org/>, June 2020)
- OCI's standards
 - Image-spec
 - Runtime-spec
 - Distribution-spec
- OCI compliant runtimes
 - containerd (used by K8s)
 - lxc
 - Runc
 - ...






Containers vs. Images

- ❑ **Image:** Imperative description (DOCKERFILE) of a collection of filesystem layers – read-only
- ❑ **Container:** Instance of an image – runtime adds a top writable layer
- ❑ The major difference between a container and an image is the top writable layer
- ❑ You can have many running containers of the same image
- ❑ Changes to the top writable layer like writing new files, changing existing ones are ephemeral. (When container exits, changes are gone → See Volumes to persist data)

<https://docs.docker.com/storage/storagedriver/>



Containers vs. Images




Image



Container
(based on ubuntu:15.04 image)

Container




Container
(based on ubuntu:15.04 image)

Corresponding Dockerfile

```
FROM ubuntu:15.04
COPY ./app
RUN make /app
CMD python /app/app.py
```

- Instantiation of image
- Runtime adds a top writable layer (ephemeral)

<https://docs.docker.com/storage/storagedriver/>




How does Docker Achieve Isolation?

1. **cgroups** (known from the lecture before)
2. **namespaces** (known from the lecture before)
3. **stackable images and copy on write**
 - Docker uses a copy-on-write mechanism when deriving new images from existing ones
 - Docker only keeps track of changes between this image and our container. Docker also pulls the corresponding layer only once.
 - All changes are organized in so-called layers, where only the uppermost layer is writable.
4. **virtual network bridges**
 - Enables communication between hosts or even not
 - Restricts accessibility

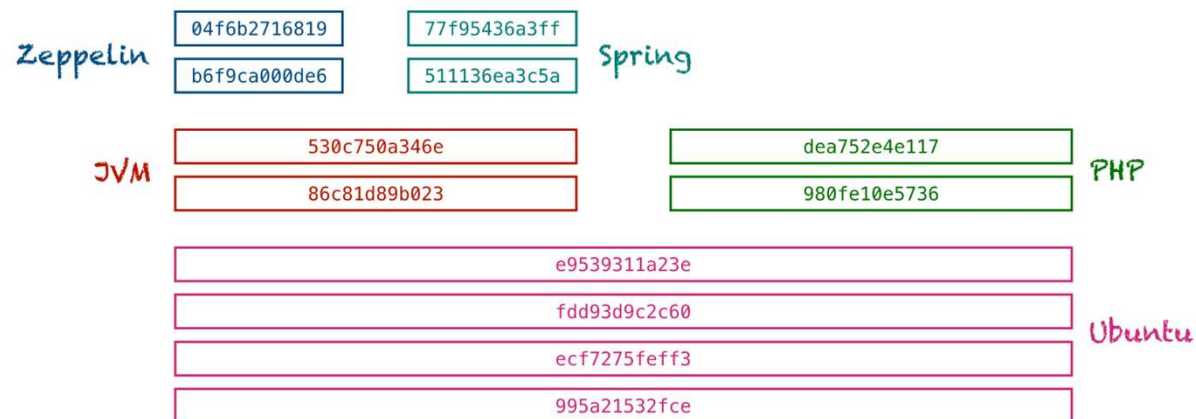
<https://blog.codecentric.de/en/2019/06/docker-demystified/>






Stackable Images and Copy-on-Write

- ❑ Layers are only stored once (Ubuntu layers are stored once and used by PHP, JVM and the other images)
- ❑ If PHP changes e.g. the top layer of Ubuntu stack (e9539..), docker copies this layer, makes changes and therefore stores it as a new layer.
- ❑ Starting a container means: putting all layers in an isolated filesystem namespace and add a top writable layer





<https://blog.codecentric.de/en/2019/06/docker-demystified/>

Networking – how can Isolated Containers Communicate?




- ❑ Each container gets an IP and is part of a network
- ❑ Docker offers different network drivers:
 - bridge: default driver – enables communication within the docker host between containers within the same bridge network
 - Automatically generated and called *bridge* and also used by default
 - Do not use default bridge network in production!
 - Use user-defined custom bridge networks
 - !!! Default bridge does not support DNS resolution
 - host: uses the docker host's networking directly
 - overlay: connects multiple daemons (some sort of a cluster)
 - none: disable networking for this container




Storage in Docker

“External” storage – why?

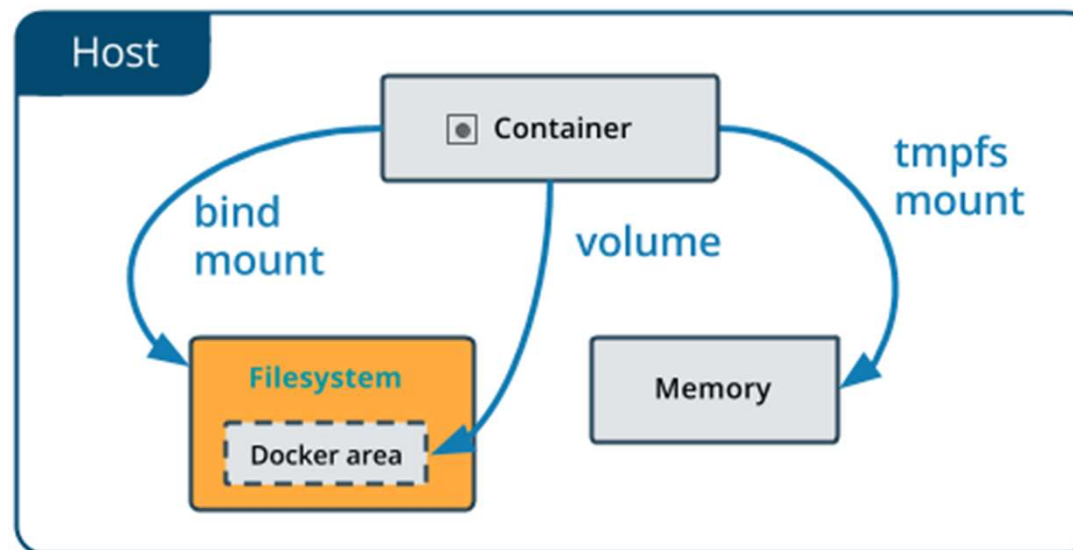


- ❑ Data written in the top writable layer only persists as long as the container is running (ephemeral data)
- ❑ Top writable layer is tightly coupled to the host, moving data is not that easy
- ❑ When writing into the top writable layer, you need a storage driver, which consumes extra resources and reduces performance




Storage in Docker

- We discuss two types of permanent storage
 - Volumes (managed by docker host – special location within host's filesystem, preferred option)
 - Bind mounts (arbitrary folder on docker host, mounted in a container)
Can be security critical when sharing config or sensitive data.



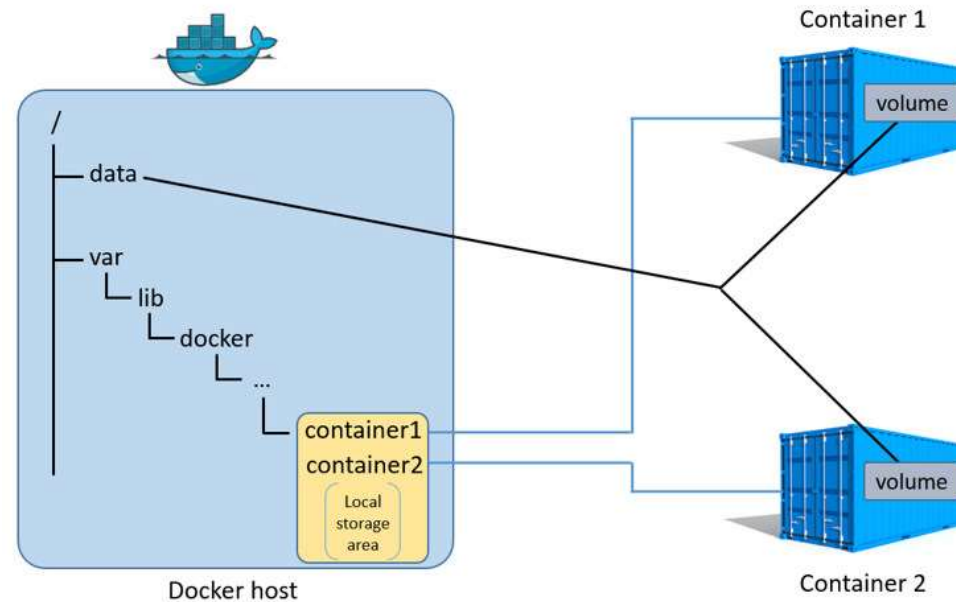
<https://docs.docker.com/storage/>



Data Volumes (preferred way)


- ❑ Data volumes are managed by docker engine
- ❑ A data volume is a directory or file in the **host's filesystem** that is mounted directly into a container
- ❑ When a container is deleted, any data written to the container that is not stored in a *data volume* is deleted along with the container
- ❑ You can mount any number of data volumes into a container
- ❑ Multiple containers can also share one or more data volumes

Data Volumes



- Each container has its own folder on the system controlled by docker (/var/lib/docker/containers/*)
- Volumes are stored under a separate folder (/var/lib/docker/volume)

Bind Mounts



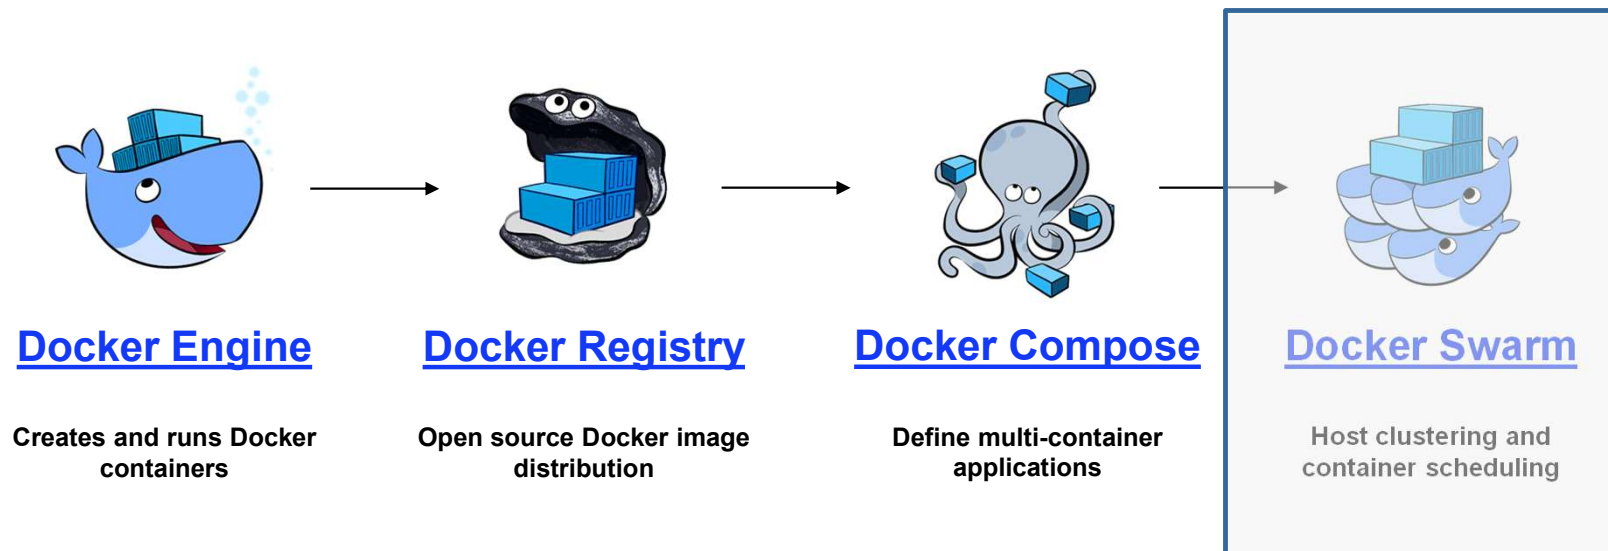

- ❑ Managed by the user (!!)
- ❑ Read-only access might be a good choice in many cases
- ❑ Sharing configuration files might be beneficial (they normally do not reside within docker's filesystem)



- ❑ **BUT:** Whenever possible –
 Use Volumes instead

<https://blogs.sap.com/2020/03/10/understanding-containers-part-05-shared-files-between-the-host-and-containers/>




(Part of) The Docker Stack






Docker Engine

Overall Docker Architecture





- ❑ Daemon: daemon of the server process to manage containers
- ❑ Client: user client to (remotely) control the daemon
- ❑ Registry: platform for sharing and managing images



<https://docs.docker.com/engine/understanding-docker/>

Engine – Overview

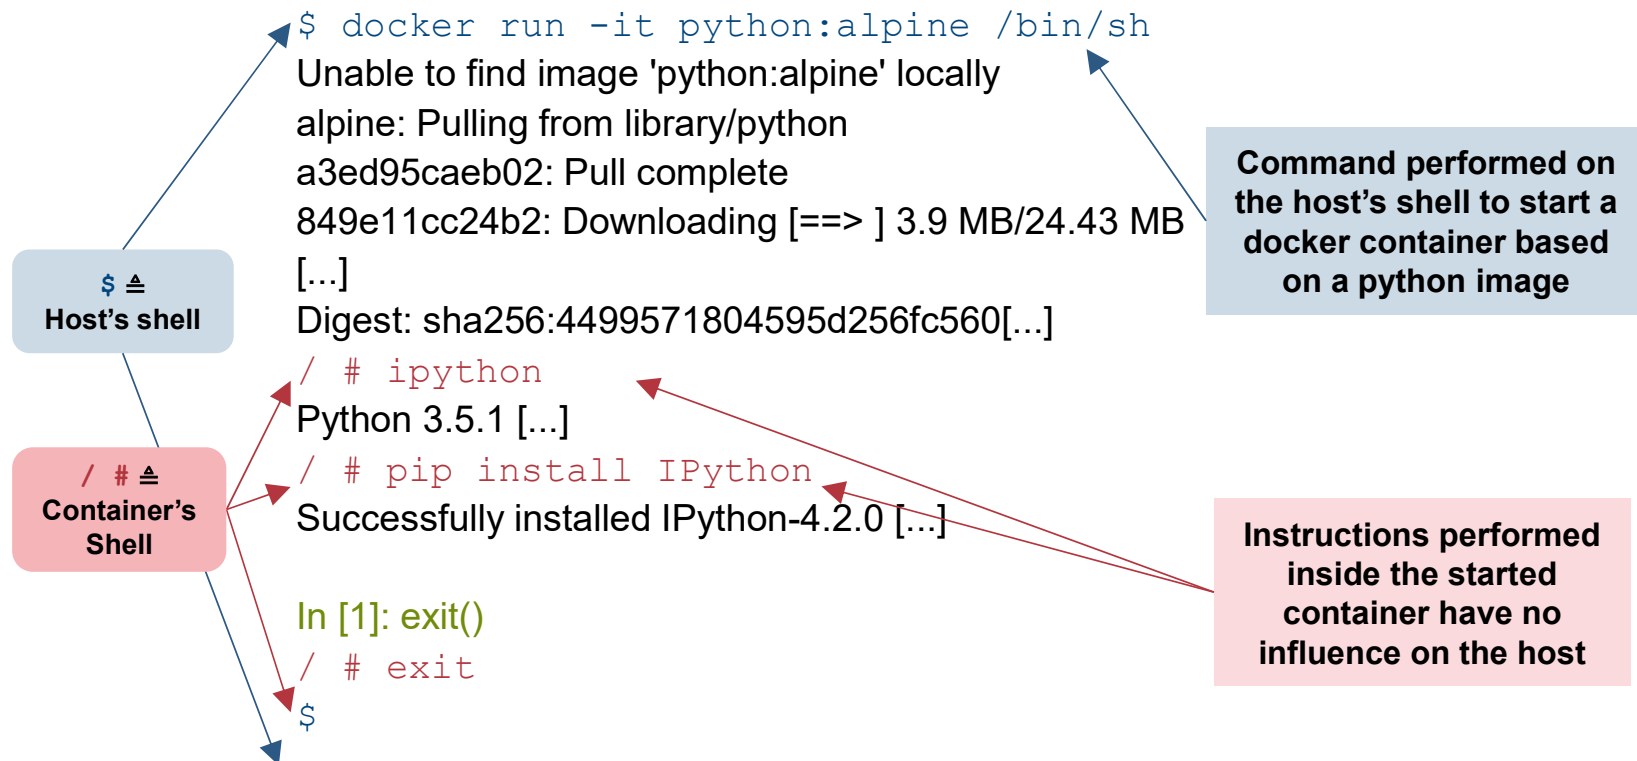




Consists of three parts:

- ❑ **Server:** Daemon of the server process to manage the containers
- ❑ **Client:** Client to (remotely) control the daemon
- ❑ **REST API**

<https://docs.docker.com/get-started/overview/>

Engine – Example





Engine – Example

□ What happened here?


- We created a *container* with its own:
 - filesystem (based initially on a `python` image)
 - network stack
 - process space
- We started a shell process (no `init`, no `systemd`, no problem)
- We installed IPython with `pip`

□ What did not happen here?

- We did not make a full copy of the `python` image
- We did not modify the `python` image itself
- We did not affect any other container (currently using this image or any other image)

□ List of comands and their explanation: Docker CLI





Engine – Dockerizing applications

- Running an application based on an image needs a single command:


```
$ docker run <IMAGE>
```

- Let's echo hello world inside a container

```
$ docker run ubuntu /bin/echo 'Hello world'
Hello world
```

- Run an interactive container

```
$ docker run -t -i ubuntu /bin/bash
root@af8bae53bdd3:/#
root@af8bae53bdd3:/# ls
bin boot dev etc home lib lib64 media mnt opt proc
root run sbin srv sys tmp usr var
root@af8bae53bdd3:/# exit
```



Engine – Working with images


□ Managing local images

\$	images	List images
	rmi	Remove one or more images
	tag	Tag an image into a repository
	inspect	Return low-level information on an image

□ Working with an image registry

\$	pull	Pull an image from a registry
	push	Push an image to a registry
	search	Search the Docker Hub for images

Engine – Working with containers




❑ Managing containers

\$	ps	List containers
	run	Run a command in a new container
	start	Start one or more stopped containers
	stop	Stop a running container
	commit	Create a new image from a container
	rm	Remove one or more containers

❑ Inspecting/debugging containers

\$	diff	Inspect changes on a container's filesystem
	inspect	Return low-level information on a container
	logs	Fetch the logs of a container
	stats	Display a resource usage statistics
	top	Display running processes of a container



Engine – Creating a new image

- Add something to an existing image **training/sinatra**

```
$ docker run -t -i training/sinatra /bin/bash
root@0b2616b0e5a8:/#
root@0b2616b0e5a8:/# gem install json
root@0b2616b0e5a8:/#exit
```


- Save the container to a new image

```
$ docker commit -m "Added json gem" -a "Kate Smith" \
0b2616b0e5a8 ouruser/sinatra:v2
4f177bd27a9ff0f6dc2a830403925b5360bfe0b93d476f7fc3231
110e7f71b1
```

- Use your new image

```
$ docker run -t -i ouruser/sinatra:v2 /bin/bash
root@78e82f680994:/#
```


Engine – Creating a new image



- ❑ Using the `docker commit` command is a pretty simple way of extending an image
- ❑ but it's a bit cumbersome and it's not easy to share a development process for images amongst a team
- ❑ Repeatable way to create images?!

→ **Dockerfiles**


Dockerfile



- ❑ Skeleton for an Image
- ❑ Contains all necessary instructions to generate an image
- ❑ Makes image creation reproducible
- ❑ Instructions, e.g., **RUN** are:
 - single line statements and contain a key word
 - not case sensitive
 - always create a new layer

```
# This is a comment
FROM ubuntu:14.04
MAINTAINER Kate Smith <ksmith@example.com>
RUN apt-get update && apt-get install -y
    ruby ruby-dev
RUN gem install sinatra
```

Dockerfile – Example



Extend base image



```
1 FROM buildpack-deps:jessie-scm
```

Run a command

*Here: install dependencies
from the package repository*



```
2  
3 # gcc for cgo  
4 RUN apt-get update && apt-get install -y --no-install-recommends \  
5     g++ \  
6     gcc \  
7     libc6-dev \  
8     make \  
9     && rm -rf /var/lib/apt/lists/*
```

Set env variables



```
10  
11 ENV GOLANG_VERSION 1.5.3  
12 ENV GOLANG_DOWNLOAD_URL https://golang.org/dl/go$GOLANG_VERSION.linux-amd64.tar.gz  
13 ENV GOLANG_DOWNLOAD_SHA256 43afe0c5017e502630b1aea4d44b8a7f059bf60d7f29dfd58db454d4e0ae53
```

Run a command

Here: download Go runtime



```
14  
15 RUN curl -fsSL "$GOLANG_DOWNLOAD_URL" -o golang.tar.gz \  
16     && echo "$GOLANG_DOWNLOAD_SHA256 golang.tar.gz" | sha256sum -c - \  
17     && tar -C /usr/local -xzf golang.tar.gz \  
18     && rm golang.tar.gz
```

Set default workdir



```
19  
20 ENV GOPATH /go  
21 ENV PATH $GOPATH/bin:/usr/local/go/bin:$PATH  
22  
23 RUN mkdir -p "$GOPATH/src" "$GOPATH/bin" && chmod -R 777 "$GOPATH"
```

Copy from host to


image



```
24 WORKDIR $GOPATH  
25  
26 COPY go-wrapper /usr/local/bin/
```



Dockerfile – Workflow




1. Define image in a **Dockerfile**
2. Use **docker build** to create a new image based on the Dockerfile

```
$ docker build .  
docker build -t ouruser/sinatra:v2 .
```

Image name

Context for building the Dockerfile,
directory might include artifacts for
building the image

3. Create a new container from your image via **docker run**



(Important) Dockerfile instructions

\$ **FROM** -- sets the Base Image for subsequent instructions (mandatory).


RUN -- will execute any commands in a new layer on top of the current image and commit the results.

CMD -- The main purpose of a CMD is to provide defaults for an executing container (only one allowed).

EXPOSE -- informs Docker that the container listens on the specified network ports at runtime.

ENV -- sets an environment variable

ADD or COPY -- copies files, directories or remote file URLs to the filesystem of the container



(Important) Dockerfile instructions


\$

ENTRYPOINT -- allows you to configure a container that will run as an executable – first command which is executed. Also see **CMD**.

VOLUME -- creates a mount point with the specified name and marks it as holding externally mounted volumes from native host or other containers.

WORKDIR -- sets the working directory for any **RUN**, **CMD**, **ENTRYPOINT**, **COPY** and **ADD** instructions.

ONBUILD -- adds to the image a trigger instruction to be executed at a later time, when the image is used as the base for another build, e.g., an application runtime container.



Best practices for images (1/2)

❑ **Containers should be ephemeral**

- can be stopped and destroyed and a new one built and put in place with an absolute minimum of set-up and configuration

❑ **Avoid installing unnecessary packages**


- to reduce complexity, dependencies, file sizes, and build times, you should avoid installing extra or unnecessary packages

❑ **Domain-driven assignment of processes to containers**

- Decoupling applications into multiple containers makes it much easier to scale horizontally and reuse containers. If that service depends on another service, make use of container linking


❑ **Minimize the number of layers**

- find the balance between readability of the Dockerfile and minimizing the number of layers it uses.




Best practices for images (2/2)

- ❑ **Keep your images as small as possible** (faster pull times, loading into memory, starting containers)
 - Start with the right base image (only the stuff you need)
 - Use multi stage builds (next slide)
- ❑ **Share layers as much as possible**
 - If you have a common basis, define your own base images
 - docker only pulls the layers once – they are cached afterwards
- ❑ **Tag your images with meaningful tags**
- ❑ **Do not store application data in the top writable layer of your container** (you know this data is ephemeral)



Best practices for Dockerfiles

- ❑ Each instruction creates one layer
 - Define the minimum of needed layers
- ❑ Use a .dockerignore file (works as .gitignore)
 - Ignores all specified files within the working directory
- ❑ Multi Stage Builds
 - Multiple bases
 - Select, copy, alter files from one stage to others
 - Name your build stages
 - You can use external images as stages
 - Only the layers of the last stage are included in the image

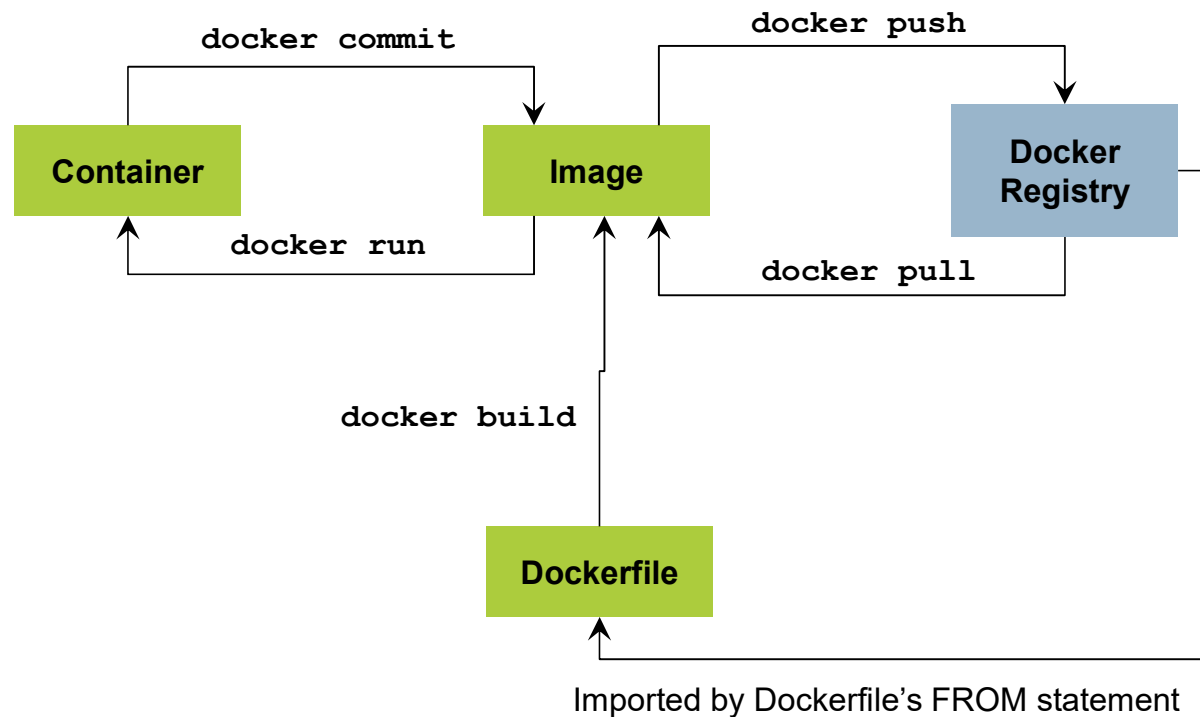




Multi-Stage Builds

```
1 # base image - builder stage
2 FROM openjdk:11.0.7-jdk AS builder
3 # Environment Variable
4 ENV APP_HOME=/root/dev/beverage
5 # Working directory
6 WORKDIR $APP_HOME
7 # Copy all the stuff (easiest way)
8 COPY . $APP_HOME
9 # Run the build
10 RUN ./gradlew build
11
12 # base image for the final image (java runtime environment is sufficient)
13 FROM openjdk:11.0.7-jre
14 # specifying work directory
15 WORKDIR /root/
16 # only copy the fat jar, which includes all dependencies (only a java runtime environment is needed to run it)
17 COPY --from=builder /root/dev/beverage/build/libs/beverage-all.jar .
18 # Run it
19 CMD ["java", "-jar", "beverage-all.jar"]
```

- ❑ Builder stage (not included in the image – only the last stage is included – beginning at last FROM statement)
- ❑ “Image stage” – All commands here result in a single layer
- ❑ Access to the builder stage and copying of the relevant file


Docker Process (so far)





Docker Registry


Registry



- ❑ Often, containers are based on or reuse existing images
- ❑ The Registry is a stateless, highly scalable server side application that stores and distributes Docker images
- ❑ Stores the layers and the description of how they make up an image
- ❑ Can be hosted locally to own the images pipeline
- ❑ Most users will be satisfied with Docker's public instance


→ **Public central registry: Docker Hub**

Docker Hub







- ❑ Centralized resource for container image discovery, distribution and change management, user and team collaboration, and workflow automation throughout the development pipeline
- **Image Repositories:** Find, manage, and push and pull images from community, official, and private image libraries
- **Automated Builds:** Automatically create new images when you make changes to a source repository
- **Organizations:** Create work groups to manage user access to image repositories



Docker Hub



- ❑ Docker Hub hosts public and private Docker images
- ❑ Docker Hub will be used for looking up missing local images or when using `docker pull`
- ❑ Includes official and community-maintained images


The screenshot shows the Docker Hub interface with a dark header containing 'Explore', 'Help', a search bar, and 'Sign up'/'Log In' buttons. The main section is titled 'Explore Official Repositories' and displays a table of official Docker images.

Repository	Stars	Pulls	Details
 nginx official	2.7K STARS	10M+ PULLS	DETAILS
 busybox official	625 STARS	10M+ PULLS	DETAILS
 ubuntu official	3.7K STARS	10M+ PULLS	DETAILS
 swarm official	296 STARS	10M+ PULLS	DETAILS

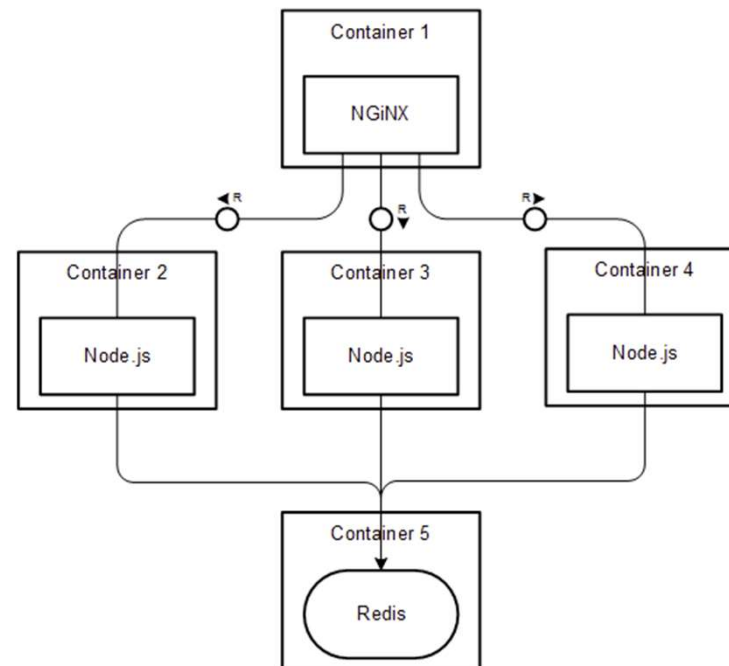


Docker Compose

Compose – The whys and wherefores




- Typically, an application is not one service but an orchestration of multiple smaller, isolated service units



<http://anandmanisankar.com/posts/docker-container-nginx-node-redis-example/>


Compose – The whys and wherefores




Multi-container apps are a hassle!

1. Build images from Dockerfiles
2. Pull existing images from the Docker Hub
3. Create Containers
4. Start/Stop Containers
5. Stream container logs

Compose – The whys and wherefores



Multi-container apps are a hassle!




```
$ docker pull redis:latest

docker build -t web .


docker run -d --name=db redis:latest
redis-server --appendonly yes

docker run -d --name=web -p 5000:5000 -v
`pwd`:./code web python app.py
```

Compose – The whys and wherefores




Multi-container apps are a hassle!




```
$ docker pull ...  
  docker pull ...  
  docker build ...  
  docker build ...  
  
  docker run ...  
  docker run ...  
  docker run ...  
  docker run ...
```


Compose – The whys and wherefores



Ideally, we want to get an app running in one command!




Compose – Overview



- ❑ Compose is a tool for defining and running multi-container Docker applications
 - ❑ Recreate a microservices architecture on development and production machines
 - ❑ All of that can be done by Compose in the scope of a single host
- For multi-host deployment, you should use more advanced solutions, like Apache Mesos or a complete Google Kubernetes architecture

Compose – Workflow



1. Define each service in a **Dockerfile**
2. Define the services and their relation to each other in the **docker-compose.yml** file
3. Use **docker-compose up** to start the system

Write your dockerfile(s)

```
WORKDIR /code
ADD requirements.txt
/code/
RUN pip install -r
requirements.txt
ADD . /code
CMD python app.py
```

Write your compose.yml file


```
web:
  build: .
  links:
    - db
  ports:
    - "8000:8000"
db:
  image: postgres
```

Run your app

```
$ docker-compose up
```

<https://docs.docker.com/compose/>

Compose – Orchestration




- ❑ The Compose file is a YAML file defining services, networks and volumes
- ❑ Configuration will be applied to each container started for that service, much like passing command-line parameters to **docker run**
- ❑ Likewise, network and volume definitions are analogous to **docker network create** and **docker volume create**
- ❑ Options specified in the Dockerfile (e.g., CMD, EXPOSE) are respected by default – no need to specify them again

Compose File Reference – A great docu!!

<https://docs.docker.com/compose/compose-file>


Compose – Example



- ❑ Two services: `web` and `redis`
- ❑ `web` is built from the `Dockerfile` in the current directory
- ❑ Forwards the exposed port 5000 on the container to port 5000 on the host machine
- ❑ Mounts the current directory on the docker host (!!) to `/code` inside the container
- ❑ `depends_on` means, that `web` is started after `redis` is ready (starting order)
- ❑ Execute **`docker-compose up [-d]`**
 - After that the `web` container should be accessible at `http://localhost:5000`

```
services:
  web:
    build: .
    ports:
      - "5000:5000"
    volumes:
      - ./code
    depends_on:
      - redis
  redis:
    image: redis
```

`docker-compose.yml`

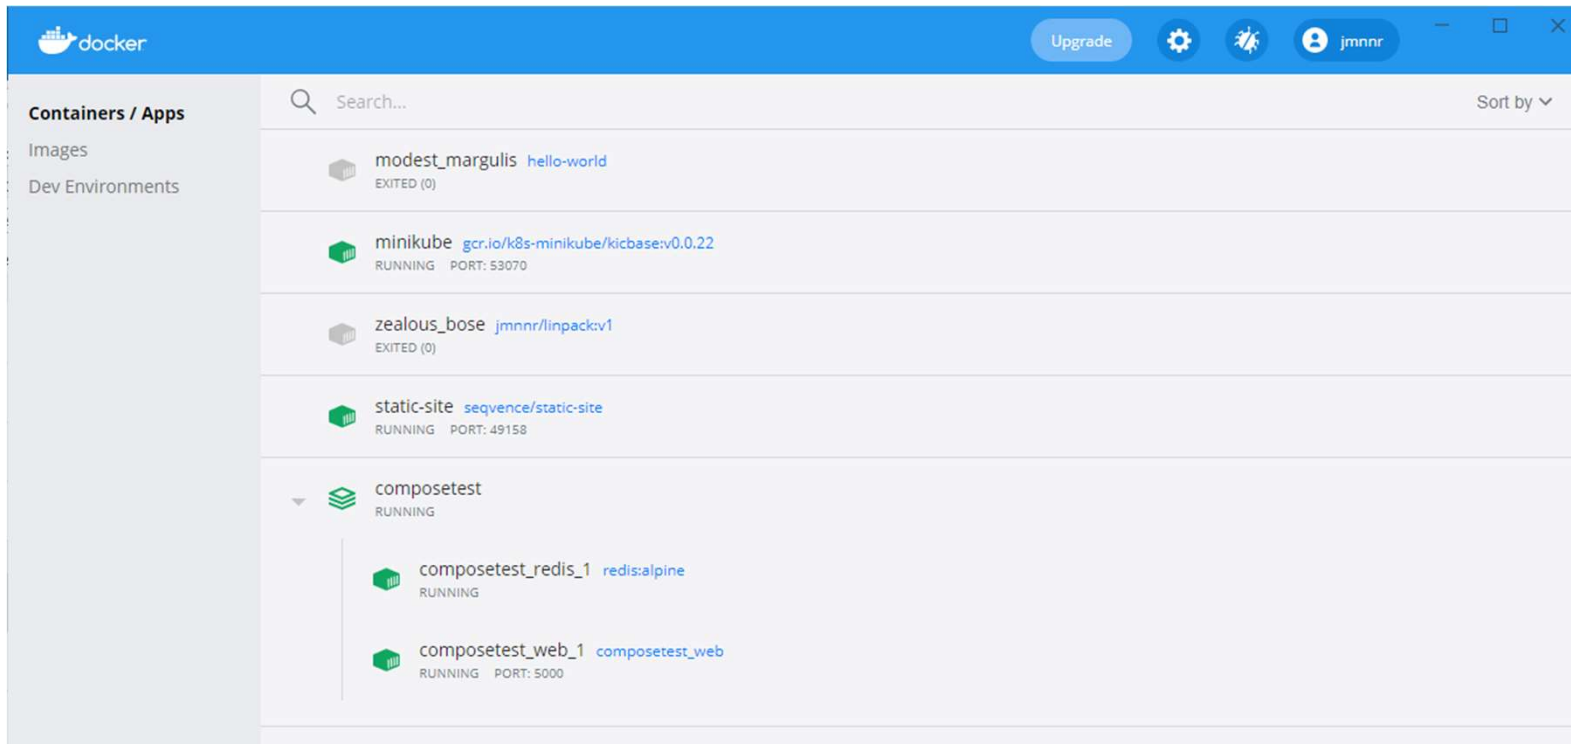



Compose – Commands


- ❑ **docker-compose** commands are a subset of docker counterparts but affect the whole multi-container architecture defined in docker-compose.yml

```
$ build -- Build or rebuild services
logs -- View output from containers
port -- Print the public port for a port binding
ps -- List containers
rm -- Remove stopped containers
run -- Run a one-off command
scale -- Set number of containers for a service
start -- Start services
stop -- Stop services
up -- Create and start containers
```

Docker Dashboard




- ❑ Benefits for development:
 - Running and terminated containers are shown in a nice view
 - Composed applications are grouped



Docker – conclusive remarks

- ❑ State of the art development and deployment of applications via containers
- ❑ Easy to use Dashboard for Windows and Mac to interact with images and container (somehow an UI for the docker CLI):
<https://docs.docker.com/desktop/dashboard/>
- ❑ Docker documentation is really great, worth reading and included in 2020/2021 a lot of information about CI/CD, building cloud native apps and deploying your containers to hosted K8s cluster or managed container services (e.g. AWS ECS)
<https://docs.docker.com/language/java/>

Related Technologies



Kubernetes is an open-source system for automating deployment, operations, and scaling of containerized applications. (Open sourced by Google)



Apache Mesos abstracts CPU, memory, storage, and other compute resources, enabling fault-tolerant and elastic distributed systems to easily be built and run effectively. (University of Berkeley)



Nomad is a tool for managing a cluster of machines and running applications on them. Nomad abstracts away machines and the location of applications, and instead enables users to declare abstract workloads. (HashiCorp)