


# DSG-SOA-M 2024: - Kubernetes -



Dr. Andreas Schönberger,  
Johannes Manner

Distributed Systems Group  
Faculty Information Systems and Applied Computer Science  
University of Bamberg





# What is Kubernetes?

## □ What is Kubernetes?

“Kubernetes (K8s) is an open-source system for automating deployment, scaling, and management of containerized applications.”


(<https://kubernetes.io/>)

## □ Why is K8s so developer friendly?

“Containerization [and therefore K8s as well] transforms the data center of being machine-oriented to being application oriented.” [BGOBW2016, p.5]

*“Kubernetes was built to radically change the way that applications are built and deployed in the cloud. Fundamentally, it was designed to give developers more velocity, efficiency, and agility.” [HBB2017, p.11]*





# Where does Kubernetes come from?

- All the systems are developed by Google Inc.


Borg*	Omega*	Kubernetes
<ul style="list-style-type: none"><li>• Built to manage long-running services and batch jobs</li><li>• Aim was to increase utilization</li><li>• Container support and isolation made sharing of resource possible</li><li>• Central component: Borg Master</li><li>• Remains the primary used container-management system within Google Inc. (?)</li></ul>	<ul style="list-style-type: none"><li>• Offspring of Borg</li><li>• Consolidation of Borg's ecosystem</li><li>• Store state in a shared persistent store</li><li>• No central component: Decoupled the Borg master's functionality</li><li>• Established improvements are folded into Borg</li></ul>	<ul style="list-style-type: none"><li>• Open source</li><li>• Designed by Google, maintained by the cloud native computing foundation (CNCF)</li><li>• Like Omega, has a shared persistent store</li><li>• Developed with a focus on external developers (public cloud provider)</li><li>• "its main design goal is to make it easy to deploy and manage complex distributed systems"</li></ul>

[BGOBW2016 ,p.3]

<https://static.googleusercontent.com/media/research.google.com/de/pubs/archive/44843.pdf>


\* Systems are closed source





# Some statistics . . .


- Survey of 400 IT and security professionals (2020)




K8s is **THE** tool to orchestrate containers

Managed K8s services become even more important compared to Elastic Container Service (ECS) or Docker Swarm

What do you use to orchestrate your containers? (pick all that apply)



Figures from <https://www.stackrox.com/kubernetes-adoption-security-and-market-share-for-containers/>




# The Success of K8s is Container Success

- ❑ Utilization -> which means in the end \$\$\$
- ❑ Containers provide isolation and dependency minimization  
-> only a few OS versions are needed
- ❑ Developers and operators do not worry about OS details
- ❑ Roll out of new hardware and OS is much easier since the container hides all details and the applications are not affected
- ❑ Self-healing of systems is much easier: If a container fails during execution or is unresponsive, the container management system spawns a new one and deletes the other
- ❑ Different pieces of a service can be developed independently from each other  
→ and K8s gives the answer to professional operation of containers


cf. Docker /  
Container  
lectures

# Agenda – The Basic Concepts



- ❑ Benefits of K8s
- ❑ Minikube
- ❑ K8s core components (K8s CC)
  - Master and Node
  - Deployment and App
  - Pods
  - Service
  - Labels and Annotations
  - Replica and DaemonSets
  - Data Integration and StatefulSets
- ❑ Outlook
- ❑ Case Study - CatService

cf. sources on  
literature slide




# What is the Benefit of K8s?

- ❑ Abstraction of infrastructure
  - Operation is much easier (abstraction of HW and OS details)
  - Easier portability due to higher level API (containers)  
(to achieve this avoid cloud managed services like DynamoDb etc.)
- ❑ Resource Efficiency
  - Cost of running a server (i.e. utilization, complexity)
  - Human resources to manage and setup infrastructure and systems
  - Containers are more lightweight than VMs or bare metal
- ❑ Development Speed
  - Iterative improvement – new services – zero downtime
  - Declarative configuration (what and not how) – self healing system (reliability)
- ❑ Supports scaling development
  - Decoupled architectures – independent scaling of parts of your system
  - Each team is responsible for a single microservice

[HBB2017, chapter 1]



# Getting hands dirty - Minikube



- ❑ Local development environment to run K8s
- ❑ Runs only a single VM – doesn't provide reliability property and other features compared to a distributed cluster
- ❑ Cool for experimentation and learning 😊
- ❑ Look at the setup page – also additional information about the concepts is provided there

<https://github.com/johannes-manner/k8s-soa-example>





**Take sometimes  
minutes for startup  
Keep calm 😊**

<https://minikube.sigs.k8s.io/docs/start/>




# K8s CC – Master and Node



**Kubernetes Cluster**

[Adapted from https://kubernetes.io/docs/tutorials/kubernetes-basics/](https://kubernetes.io/docs/tutorials/kubernetes-basics/)

# K8s CC – Master and Node



Kubernetes cluster – two types of resources:

## ❑ Master

- Managing the cluster
- Coordinates scheduling, maintaining desired state, scaling, rolling out new updates etc.

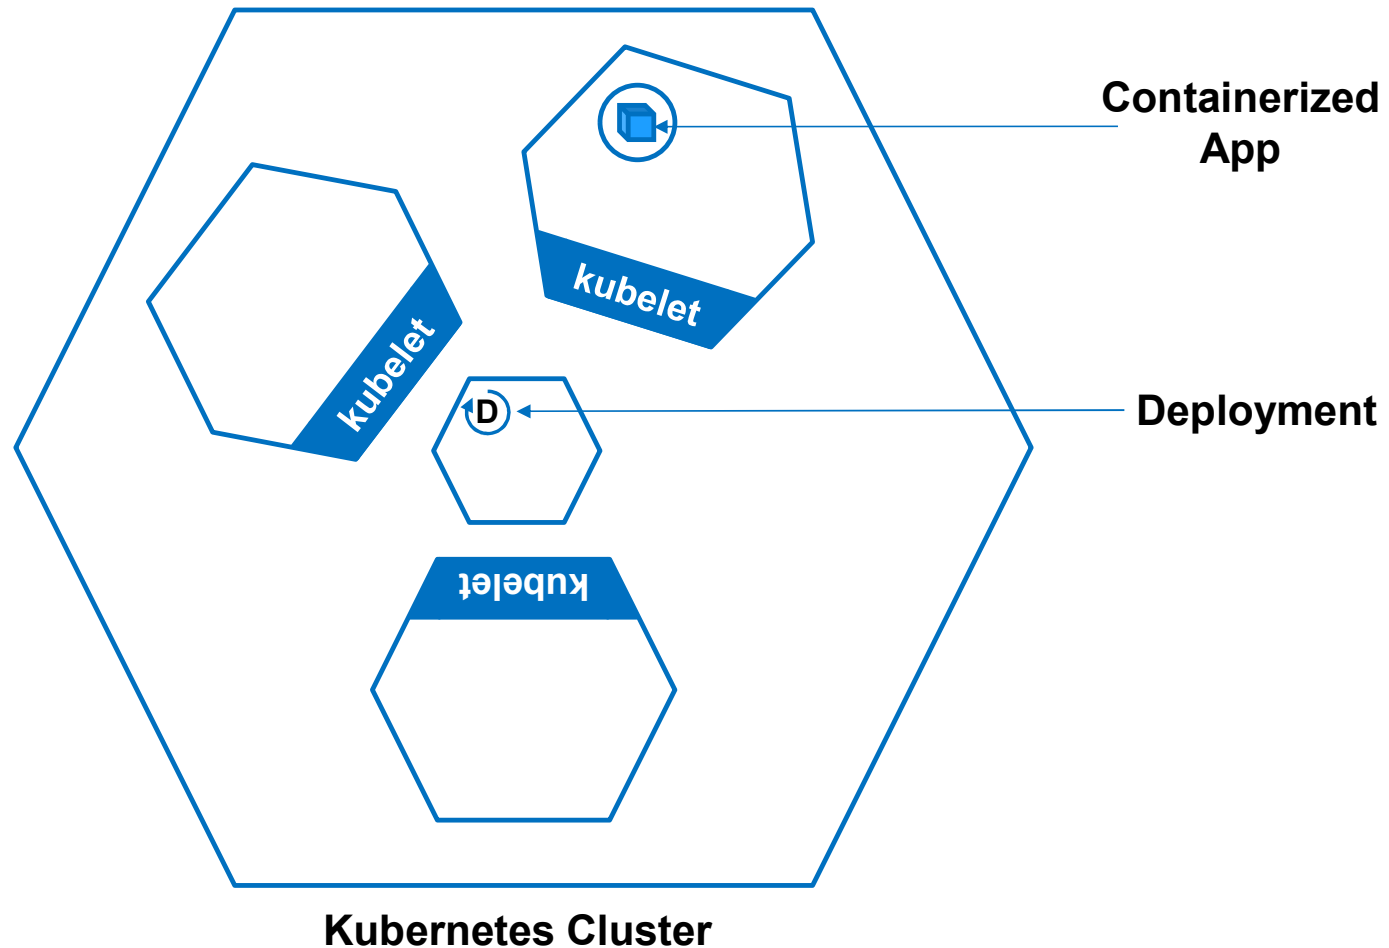

## ❑ Node

- VM or physical computer
- Serves as a worker within the cluster


## ❑ Kubelet

- Agent for managing node
- Node process to communicate with the master (Kubernetes API)

# K8s CC – Deployment and App



Adapted from <https://kubernetes.io/docs/tutorials/kubernetes-basics/>



# K8s CC – Deployment and App

## □ Deployment

- Instruction what to create and update
- After deployment creation, master schedules apps on nodes
- Deployment Controller compares the desired state (declarative config) with the observed state and takes actions if necessary
  - > self-healing if machine failures occur
- Creating and updating deployments via kubectl (K8s CLI)

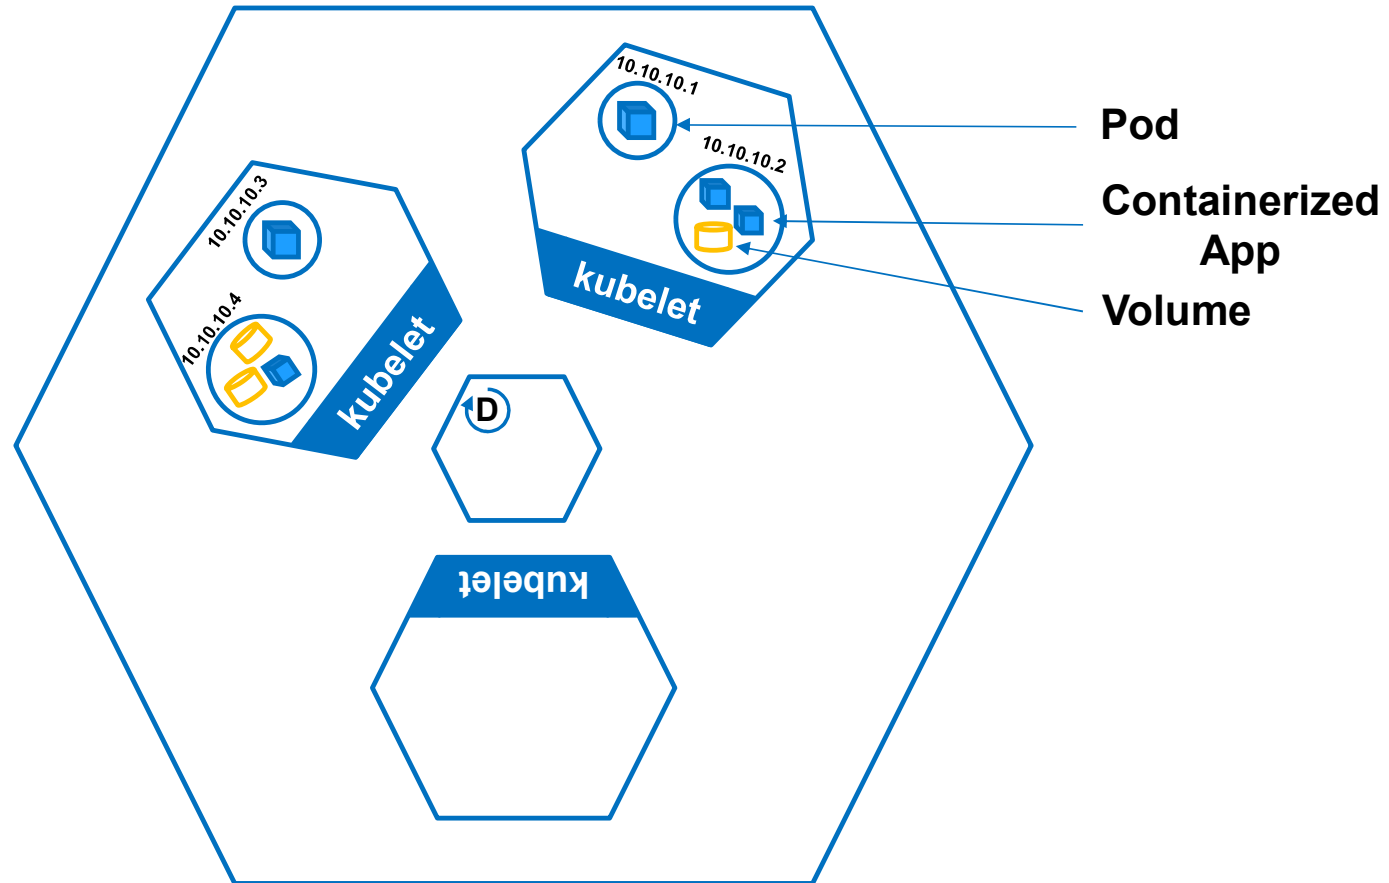

```
$ kubectl create deployment DEPLOYMENT-NAME  
    --image=ADDRESS ...
```

- Proxy to connect to the private network via `$kubectl proxy`  
(shows management information via a REST based API)

## □ App

- Has to be containerized (see Docker slides)
- Number of Replicas, ideally 3 (why?), and other settings can be specified/changed via kubectl commands


# K8s CC – Pods



**Kubernetes Cluster**


[Adapted from https://kubernetes.io/docs/tutorials/kubernetes-basics/](https://kubernetes.io/docs/tutorials/kubernetes-basics/)





## K8s CC – Pods

- ❑ **Smallest deployable unit of K8s**
- ❑ Abstraction for a group of application containers and volumes accessible via a unique IP
- ❑ Pod IP is not exposed outside the cluster  
For debugging, use port forwarding capability of K8s
- ❑ Each container in a pod runs in its own cgroup  
Process health checks are defined per container  
(executed every 10s, if failing 3 consecutive times – container restart)
- ❑ A pod does not move, once scheduled  
(solution: destroy or reschedule the pod)
- ❑ Resources are requested per container
- ❑ Lifecycle considerations: <https://dzone.com/articles/kubernetes-lifecycle-of-a-pod>



# K8s CC – Pods Declarative Configuration

## kuard-pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: kuard
spec:
  containers:
  - image: gcr.io/kuar-demo/kuard-amd64:1
    name: kuard
    ports:
    - containerPort: 8080
      name: http
      protocol: TCP
```



The port 8080 is only visible within the K8s cluster, exposing the port via `$kubectl port-forward POD-NAME SERVICE-PORT:POD-Port`, in our example `$kubectl port-forward kuard 8989:8080`  
Access the pod via `localhost:8989`

```
$kubectl apply -f kuard-pod.yaml
```

[HBB2017, p.41]



# K8s CC – Service




**Kubernetes Cluster**

Adapted from <https://kubernetes.io/docs/tutorials/kubernetes-basics/>





# K8s CC – Service



## □ Service


- Abstraction for a logical set of pods with an own IP; enable self-healing
- Set of pods is determined by a LabelSelector
- Enables applications to receive traffic
- Types to expose a service
  - Cluster IP (default): service only reachable within the cluster
  - NodePort: exposes the service externally on each selected node with the same port (only for debugging useful)
  - LoadBalancer (our recommended way): Creates a load balancer and assigns a fixed, external reachable IP to it
  - ExternalName: Exposes service with an external name (kube-dns)

# K8s CC – Labels and Selectors




Adapted from <https://kubernetes.io/docs/tutorials/kubernetes-basics/>

# K8s CC – Labels and Selectors

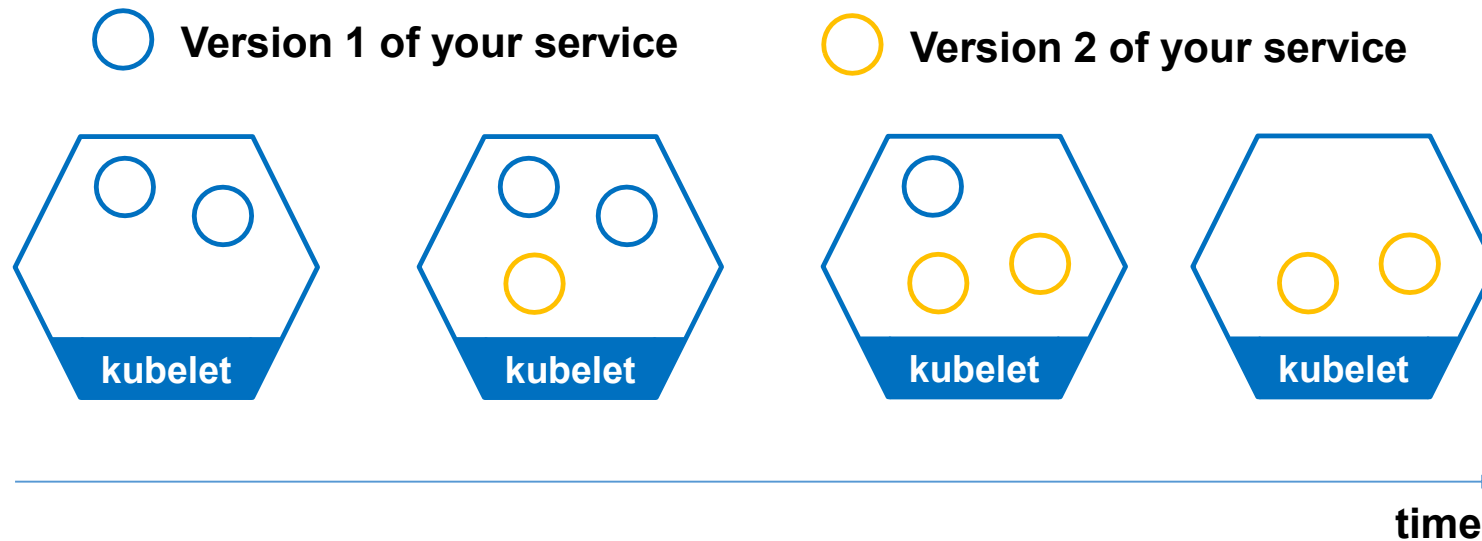


- ❑ Labels – foundation for grouping objects
  - Are key/value pairs to add metadata to your objects
  - “tier”: “frontend” – “tier”: “backend” are good examples
  - Identifying metadata - Many objects can use the same label names (grouping, viewing, operating)
- ❑ Annotations
  - Nonidentifying information for tools, libraries or as documentation
- ❑ Label Selectors
  - Group, filter a set of objects
  - Types of selectors: equality-based and set-based selectors
    - Equality-based: == and !=, e.g. environment == production  
Concatenation via “,”, e.g. a=b,c=d (means: select all objects with key/value pairs a/b and c/d)
    - Set-based : allow filtering via a set of keys and values
  - Pods of a service share the same label(s)
  - ReplicationController also uses selectors to handle pods




# K8s CC – Scaling and rolling updates


- ❑ Scaling to zero and autoscaling is supported
- ❑ Services use an integrated load balancer for request distribution
- ❑ With more than one pod, you can perform rolling updates with zero downtime



# K8s CC – Further Concepts




- ❑ Namespaces
  - Used for organizing objects in the cluster
  - Level of strong separation
- ❑ “IP for everything within the cluster”
  - Kubectl makes HTTP requests to access K8s objects
- ❑ Declarative configuration
  - Objects are represented as JSON or YAML files
  - Creating an object: `$kubectl apply -f obj.yaml`
  - Deleting an object: `$kubectl delete -f obj.yaml`
  - YAML is preferred since it is more human readable and allows adding comments



## K8s CC – ReplicaSets

- ❑ Running multiple replicas of the same container, reasons:
  - Redundancy (Avoiding outage of the service)
  - Scaling (Handling more or less requests dynamically)
  - Sharding (Partitioning of tasks in parallel)
- ❑ Treating a set of pods as logical entity
- ❑ Basis for self-healing property of K8s since pods are rescheduled under specific failure conditions
- ❑ **Replication Controller** compares desired and actual state and takes action if necessary, e.g. start /delete a pod
- ❑ Monitoring is done by labels, a ReplicaSet knows its pods via the same labels




## K8s CC – DaemonSets

- ❑ Ensures a single copy on every node in the K8s cluster
- ❑ Examples: Logging and/or monitoring agents
- ❑ Managed by a replication controller

### Difference to ReplicaSets

- ❑ Multiple replicas managed by a ReplicaSet can be placed on the same node, whereas DaemonSets schedule a single instance on each node in the cluster

# K8s CC – Self healing property - Example



```
$ kubectl apply -f cats-rs.yaml
```

Created a ReplicaSet with 3 replicas

```
$ kubectl get pods
```

Unable to connect to the server: dial tcp 192.168.99.107:8443: connectex:

Minikube crashed

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
cats-rs-jh8hm	0/1	Error	1	3m2s
cats-rs-qcstl	0/1	Error	2	3m2s
cats-rs-tgvjl	0/1	Error	1	3m2s

No Pods are ready to serve requests

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
cats-rs-jh8hm	1/1	Running	2	3m15s
cats-rs-qcstl	0/1	CrashLoopBackOff	2	3m15s
cats-rs-tgvjl	0/1	Error	2	3m15s

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
cats-rs-jh8hm	1/1	Running	2	3m28s
cats-rs-qcstl	0/1	CrashLoopBackOff	2	3m28s
cats-rs-tgvjl	0/1	CrashLoopBackOff	2	3m28s


```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
cats-rs-jh8hm	1/1	Running	2	3m37s
cats-rs-qcstl	1/1	Running	3	3m37s
cats-rs-tgvjl	1/1	Running	3	3m37s

Self healing (replication controller) of K8s starts/reschedules new pods









## K8s CC – Data Integration

- ❑ Often the most complicated step in building distributed systems in container and container orchestration environments
- ❑ Due to legacy systems and evolution of the cloud, data services are often externalized cloud services
  - > difficult/impossible to host them in a local cluster
- ❑ Two Options for data integration in K8s
  1. Externalized cloud service like DynamoDb (default for many enterprise scenarios – think about ease of use)
  2. Using storage solutions within K8s
    - a) Running Reliable Singletons
    - b) Using StatefulSets



## K8s CC – DI – 2a) Reliable Singletons


- ❑ Use K8s features like ReplicaSets, BUT not replicate the storage
- ❑ As reliable as running a single database VM or physical server, i.e. more reliable since replication controller handles failures or other outages of your singleton
  - > quite simple (no storage replication)
  - > only short downtime (for many systems acceptable)



## K8s CC – DI – 2b) StatefulSets

### Difference to ReplicaSets


- ❑ In a ReplicaSet, all pods are homogeneous (can be replaced without any action from the K8s orchestrator) – Pods in a StatefulSet have the following properties:
  - Each replica gets a unique index and name( e.g., mongodb-0, mongodb-1 etc.)
  - Creation is sequentially and blocks until the previous one is healthy and running
  - Deletion of pods is the other way around; starting from the highest index
- ❑ The service for a StatefulSet is “headless”, means that the service has no ClusterIP since all pods are unique and have its own access point




## K8s CC – DI – 2b) StatefulSets and Volumes

- ❑ Replicate the stateful layer:  
Specify a primary replica (normally pod xy-0) and add all the other database replicas to the primary replica for synchronizing state
- ❑ Persistent storage:  
Mount a volume into a specific folder of your pod (in mongoDB into /data/db)
- ❑ For each created database replica, a single volume is created too, so be careful with consistency (read the docs of the used databases)

# Outlook – what's missing right now



- ❑ Monitoring and Health Checks
- ❑ Persistent Volumes
- ❑ Jobs (third form of replicated pods)
- ❑ ConfigMaps
- ❑ Secrets
- ❑ ...



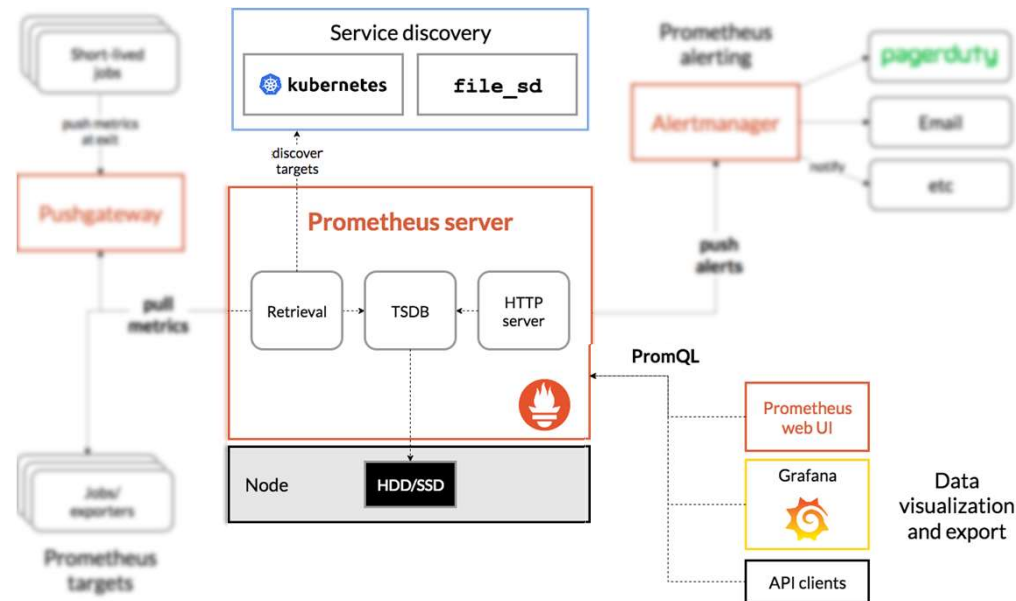

# Monitoring

- ❑ Monitoring and getting information is not that easy in a containerized application (all entities are encapsulated)
- ❑ Use Pod's container information via:
  - `$ kubectl describe pod <pod-name>`
  - `$ kubectl logs <pod-name> -c <container-name>`
- ❑ Specific information can be found in K8s docu:  
<https://kubernetes.io/docs/tasks/debug-application-cluster/>
- ❑ Tools for Monitoring: Prometheus (<https://prometheus.io/>)
  - a multi-dimensional data model with time series data identified by metric name and key/value pairs
  - no reliance on distributed storage; single server nodes are autonomous
  - multiple modes of graphing and dashboarding support


Tool info copied from: <https://prometheus.io/docs/introduction/overview/>



# Monitoring: Prometheus Architecture



- **State of the art:** Using a combination of Kubernetes for container orchestration, Prometheus for collecting metrics and Grafana for visualizing this data



# Health Checks

- ❑ K8s checks the health of your containers via “process health checks” permanently
- ❑ If container fails this check, K8s restarts it.
- ❑ Checks are application specific and included in YAML files

## Liveness

- Specify endpoint\* for liveness probes
  - Timeout, interval and threshold
  - If container does not respond in appropriate time, container is restarted (that's the default; dependent on restart policy)
- Result:** application is running properly



## Readiness

- Specify endpoint\* for readiness probes
  - Timeout, interval and threshold
  - If container does not respond in appropriate time, container is removed from service load balancers
- Result:** application is ready to serve requests

\* Endpoints can be a HTTP endpoint (e.g. for microservices) or tcpSocket (e.g. for databases), dependent on the service.




# Persistent Volumes Subsystem



- ❑ “The PersistentVolume subsystem provides an API for users and administrators that abstracts details of how storage is provided from how it is consumed”
  - PersistentVolume (PV):
    - Created by a cluster administrator
    - Piece of Storage in your cluster (handled like any other Kubernetes object)
    - Lifecycle independent of individual pods
  - PersistentVolumeClaim (PVC):
    - Request for storage by a user
    - PVCs consume PV resources
    - Claims can access specific size and access modes on a PV (decoupling)
- ❑ Control loop in master watches for PVCs and binds them to matching PVs
- ❑ Pods use claims as volumes. Per default the master looks dynamically for a matching PV. (Use label selectors)

<https://kubernetes.io/docs/concepts/storage/persistent-volumes/>

# Jobs – Ephemeral Task Executions




- ❑ Runs a specific task to completion  
No pods are running when the task is done
- ❑ Creates pods as specified (also restarts pods of jobs, which are not completed yet)
- ❑ When job completes, pods are not deleted (kept for investigation, i.e. container logs)
- ❑ Deleting the job also deletes the associated pods
- ❑ You can also define periodic jobs (cron jobs)

<https://kubernetes.io/docs/concepts/workloads/controllers/job/>

<https://kubernetes.io/docs/concepts/workloads/controllers/cron-jobs/>





# ConfigMaps and Secrets

- ❑ ConfigMaps: Set of external variables for container
- ❑ Secrets: For sensitive environment variables/data
- ❑ Three ways of using a ConfigMap
  - Filesystem: mount ConfigMap in a Pod
  - Environment variable: set values of environment variables
  - CLI argument: K8s creates CLI argument based on ConfigMap

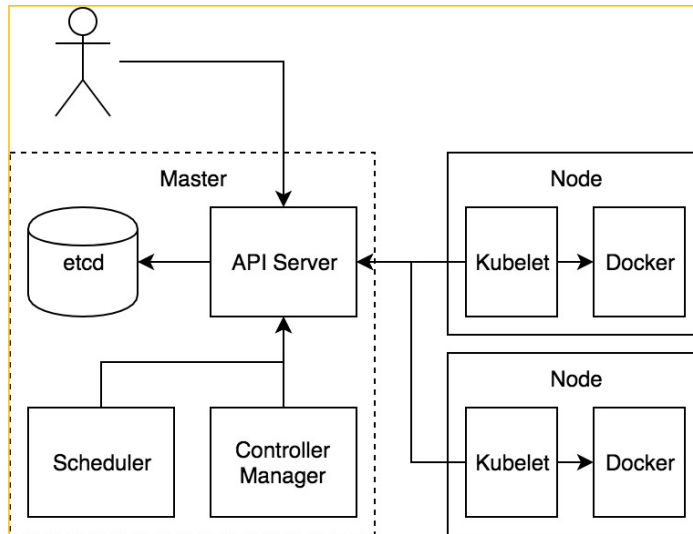

## ConfigMap YAML

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: special-config
  namespace: default
data:
  serverUri: very-best
```

## Part of Pod's YAML

```
env:
  - name: DEMO
    value: "Test it"
  - name: SERVER_URI
    valueFrom:
      configMapKeyRef:
        name: special-config # name of the config map
        key: serverUri      # key of the config property
```

# Wrap up: K8s System Architecture



## Components

**etcd:** the core state store for Kubernetes (hidden, only accessible via API Server!)

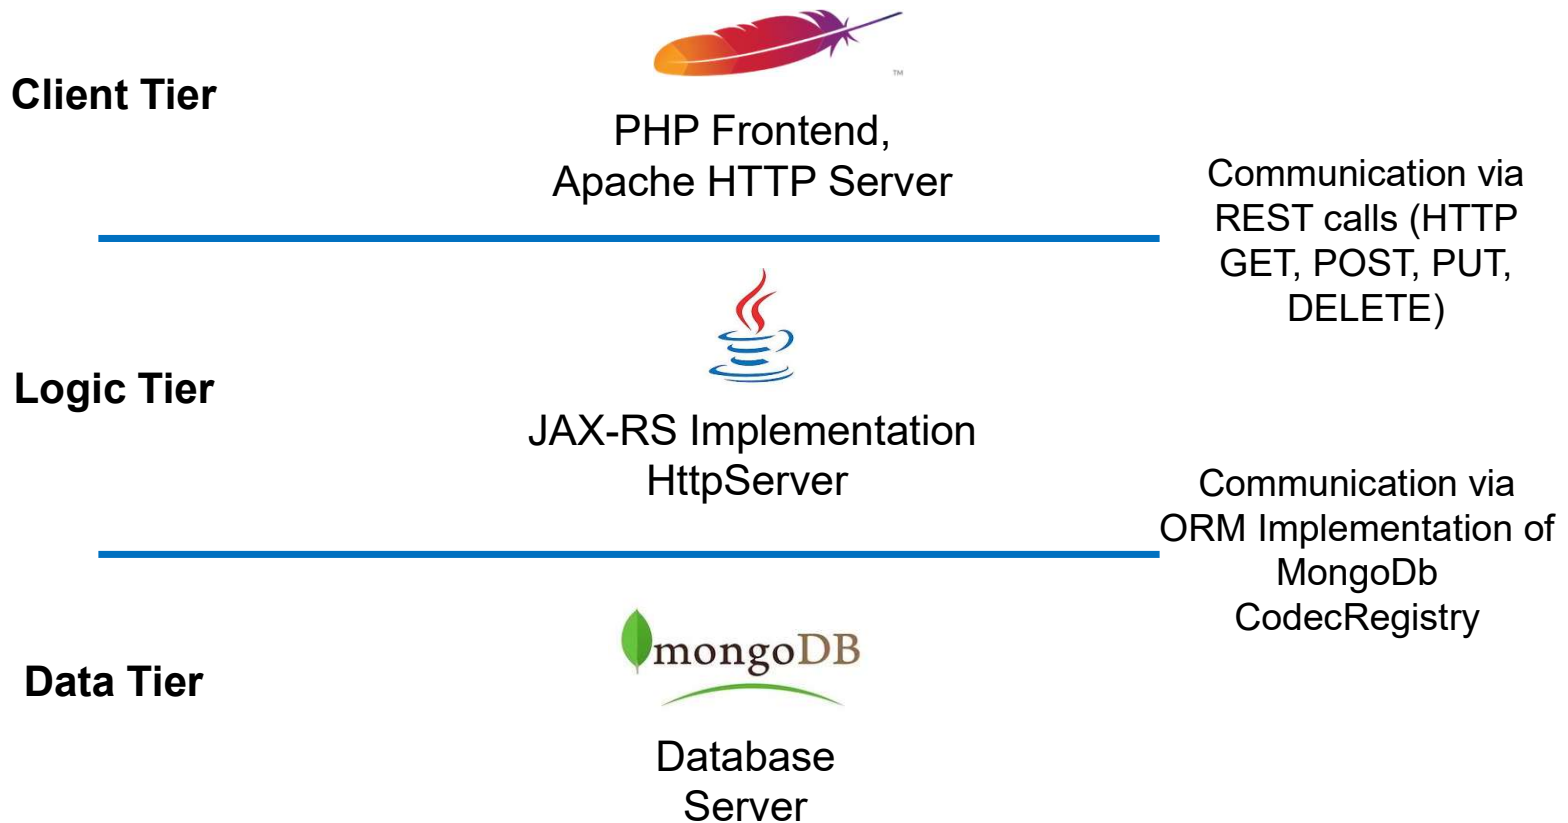

**API Server:** Simple K8s REST API  
Responsible for Authentication/ Authorization, admission controller (reject or adjust requests), API versioning.

**Scheduler/Controller Manager:** makes system work: bind and unbind Pods to nodes. Place of the reconciliation controller, etc.


**Kubelet:** K8s agent on every node

<https://blog.heptio.com/core-kubernetes-jazz-improv-over-orchestration-a7903ea92ca>

# Use Case: Cat Service - Architecture



# Use Case: Cat Service – Dashboard (1/2)



Search

+ CREATE

Overview

## Cluster

Namespaces  
Nodes  
Persistent Volumes  
Roles  
Storage Classes

Namespace

default

## Overview

## Workloads


Cron Jobs  
Daemon Sets  
Deployments  
Jobs  
Pods  
Replica Sets  
Replication Controllers  
Stateful Sets

## Discovery and Load Balancing


Ingresses  
Services

## Workloads


### Workloads Statuses




Deployments



Pods



Replica Sets



Stateful Sets

### Deployments

Name	Labels	Pods	Age	Images
cats-d-b-db	app: cats-b-db	1 / 1	3 minutes	jmnnr/soa-k8s-backend-d...
cats-d-f	app: cats-f	1 / 1	3 minutes	jmnnr/soa-k8s-frontend.l...

### Pods

Name	Node	Status	Restarts	Age
cats-d-b-db-6c9f5947f9-5f6jq	minikube	Running	0	3 minutes
mongod-0	minikube	Running	0	3 minutes
cats-d-f-d55688787-fr7lh	minikube	Running	0	3 minutes



etes  
berg

## Config and Storage

Config Maps

Persistent Volume Claims

Secrets

## Settings

About

# Replica Sets Use Case: Cat Service – Dashboard (2/2)

Name	Labels	Pods	Age	Images		
✓ cats-d-b-db-6c9f5947f9	app: cats-b-db pod-template-hash: 6c9f...	1 / 1	3 minutes	jmnrr/soa-k8s-backend-d...	☰	⋮
✓ cats-d-f-d55688787	app: cats-f pod-template-hash: d556...	1 / 1	3 minutes	jmnrr/soa-k8s-frontend:l...	☰	⋮

## Stateful Sets

Name	Labels	Pods	Age	Images		
✓ mongod	app: db	1 / 1	3 minutes	jmnrr/soa-k8s-database:l...	☰	⋮

## Discovery and Load Balancing

## Services


Name	Labels	Cluster IP	Internal endpoints	External endpoints	Age	
✓ cat-service-backend	-	10.98.74.142	cat-service-backen...	-	3 minutes	⋮
✓ mongodb-service	name: db	None	mongodb-service:...	-	3 minutes	⋮
⌚ cat-service-frontend	-	10.96.175.13	cat-service-fronten... cat-service-fronten...	-	3 minutes	⋮
✓ kubernetes	component: apiser.. provider: kubernete	10.96.0.1	kubernetes:443 TCP	-	10 minutes	⋮

## Config and Storage

## Persistent Volume Claims

Name	Status	Volume	Capacity	Access Modes	Storage Class	Age	
✓ mongodb-persi...	Bound	pvc-aa661a2e-81ea-11e9-b112-080027c88d7c	1Gi	ReadWriteOnce	-	3 minutes	⋮





## Literature – Interesting sources

- ❑ [BGOBW2016] B. Burns et al., “Borg, Omega, and Kubernetes,” ACM Queue, vol. 14, pp. 70–93, 2016
- ❑ [HBB2017] K. Hightower, B. Burns, and J. Beda,: Kubernetes: Up and Running. O’Reilly UK Ltd., 2017.  
German version: <https://katalog.ub.uni-bamberg.de/query/BV044795699>
- ❑ Nigel Poulton: The Kubernetes Book (e-Book, frequently updated, but not that mature as “up and running”)
- ❑ <https://kubernetes.io/docs/tutorials/#basics>

Second  
edition  
available

- ❑ CatService UseCase: <https://github.com/johannes-manner/k8s-soa-example>

